



collapse: Advanced and Fast Statistical Computing and Data Transformation in R

Sebastian Krantz 

Kiel Institute for the World Economy

Abstract

collapse is a large C/C++-based infrastructure package facilitating complex statistical computing, data transformation, and exploration tasks in R—at outstanding levels of performance and memory efficiency. It also implements a class-agnostic approach to R programming, supporting vector, matrix and data frame-like objects and their popular extensions (`'units'`, `'integer64'`, `'xts'`, `'tibble'`, `'data.table'`, `'sf'`, `'pdata.frame'`), enabling its seamless integration with the bulk of the R ecosystem. This article introduces the package's key components and design principles in a structured way, supported by a rich set of examples. A small benchmark demonstrates its computational performance.

Keywords: statistical computing, vectorization, data manipulation and transformation, class-agnostic programming, summary statistics, C/C++, R.

1. Introduction

collapse is a large C/C++-based R package that provides an integrated suite of statistical and data manipulation functions.¹ Core functionality includes a rich set of S3 generic (grouped, weighted) statistical functions for vectors, matrices, and data frames, which provide efficient low-level vectorizations, OpenMP multithreading, and skip missing values by default (`na.rm = TRUE`). It also provides powerful data manipulation functions, including vectorized and verbose hash-joins and fast (aggregation, recast) pivots, functions and classes for indexed (time-aware) computations on time series and panel data, recursive tools to deal with nested data, and advanced descriptive statistical tools. This functionality is powered by efficient algorithms for grouping, ordering, deduplication, and matching callable at both R and C levels. The package also includes efficient object converters, functions for memory efficient R

¹Website: <https://sebkranz.github.io/collapse/>. Linecount (v2.1.2): R: 13,785, C: 19,013, C++: 9,844. Exported namespace: 392 objects, of which 238 functions (excl. methods and shorthands), and 2 datasets.

programming, such as (grouped) transformation and math by reference, and helper functions to handle variable labels, attributes, and missing data. **collapse** is **class-agnostic**, providing statistical operations on vectors, matrices, and data frames, and seamlessly supporting popular extensions to these objects such as `'units'`, `'integer64'`, `'xts'`, `'tibble'`, `'data.table'`, `'sf'`, and `'pdata.frame'`. It is also extensively globally and interactively configurable.

Why combine all of these features in a package? The short answer is to make computations in R as flexible and powerful as possible. The more elaborate answer is to (1) facilitate complex data transformation, exploration, and computing tasks in R; (2) increase the performance and memory efficiency of R programs; and (3) create a new foundation package for statistics and data manipulation that implements many successful ideas in the R ecosystem and other programming environments in a stable, high performance, and broadly compatible manner.²

R already has a large and tested data manipulation and statistical computing ecosystem. Notably, the **tidyverse** (Wickham *et al.* 2019) provides a consistent toolkit for data manipulation in R, centered around the `'tibble'` (Müller and Wickham 2023) object and tidy data principles (Wickham 2014). **data.table** (Dowle and Srinivasan 2025) provides an enhanced high-performance data frame with parsimonious syntax. **sf** (Pebesma 2018) provides a data frame for spatial data and supporting functionality. **tsibble** (Wang *et al.* 2020) and **xts** (Ryan and Ulrich 2023) provide classes and operations for time series data, the former via an enhanced `'tibble'`, the latter through an efficient matrix-based class. Econometric packages like **plm** (Croissant and Millo 2008) and **fixest** (Bergé 2018) also provide solutions to deal with panel data and irregularity in the time dimension. Packages like **matrixStats** (Bengtsson 2025) and **Rfast** (Papadakis *et al.* 2023) offer fast statistical calculations along the rows and columns of matrices as well as faster statistical procedures. **DescTools** (Signorell 2023) provides a wide variety of descriptive statistics, including weighted versions. **survey** (Lumley 2004) offers computations on complex surveys. **labelled** (Larmarange 2023) provides tools to deal with labelled data. Packages like **tidyr** (Wickham *et al.* 2024), **purrr** (Wickham and Henry 2023) and **rapply** (Chau 2022) provide functions to deal with nested or messy data.

collapse relates to and integrates key elements from these projects. It offers **tidyverse**-like data manipulation at the speed and stability of **data.table** for any data frame-like object. It can turn any vector/matrix/data frame into a time-aware indexed series or frame and perform operations such as lagging, differencing, scaling or centering, encompassing and enhancing core manipulation functionality of **plm**, **fixest**, and **xts**. It also performs fast (grouped, weighted) statistical computations along the columns of matrix-like objects, complementing and enhancing **matrixStats** and **Rfast**. Its low-level vectorizations and workhorse algorithms are accessible at the R and C-levels, unlike **data.table**, where most vectorizations and algorithms are internal. It also supports variable labels and intelligently preserves the attributes of all objects, complementing **labelled**. It provides novel recursive tools to deal with nested data, enhancing **tidyr**, **purrr**, and **rapply**. Finally, it provides a small but consistent and powerful set of descriptive statistical tools, yielding sufficient detail for most data exploration purposes, requiring users to invoke packages like **DescTools** or **survey** only for specialized statistics.

²Such ideas include **tidyverse** syntax, vectorized aggregations (**data.table**), data transformation by reference (**pandas**), vectorized and verbose joins (**polars**, STATA: StataCorp LLC. (2023)), indexed time/panel series (**xts**, **plm**), summary statistics for panel data (STATA), variable labels (STATA), recast pivots (**reshape2**), etc...

Other programming environments such as Python and Julia by now also offer computationally very powerful libraries for tabular data such as **DataFrames.jl** (Bouchet-Valat and Kamiński 2023), **Polars** (Vink *et al.* 2023), and **Pandas** (McKinney 2010; **pandas** Development Team 2023), and supporting numerical libraries such as **Numpy** (Harris *et al.* 2020), or **StatsBase.jl** (JuliaStats 2023). In comparison with these, **collapse** offers a class-agnostic approach bridging the divide between data frames and atomic structures, has more advanced statistical capabilities,³ supports recast pivots and recursive operations on lists, variable labels, verbosity for critical operations such as joins, and is extensively globally configurable. In short, it is very useful for complex statistical workflows, rich datasets (e.g., surveys), and for integrating with different parts of the R ecosystem. On the other hand, **collapse**, for the most part, does not offer a sub-column-level parallel architecture and is thus not extremely competitive with top frameworks, including **data.table**, on aggregating billion-row datasets with few columns.⁴ Its vectorization capabilities are also limited to the statistical functions it provides and not, like **DataFrames.jl**, to any Julia function. However, as demonstrated in Section 3, vectorized statistical functions can be combined to calculate more complex statistics in a vectorized way.

The package comes with a built-in structured **documentation** facilitating its use. This includes a central **overview page** linking to all other documentation pages and a set of supplementary topic pages which briefly summarize related functionality. The names of these extra pages are collected in a global macro `.COLLAPSE_TOPICS` and can be called directly with `help()`:

```
R> .COLLAPSE_TOPICS
```

```
[1] "collapse-documentation"      "fast-statistical-functions"
[3] "fast-grouping-ordering"      "fast-data-manipulation"
[5] "quick-conversion"           "advanced-aggregation"
[7] "data-transformations"        "time-series-panel-series"
[9] "list-processing"             "summary-statistics"
[11] "recode-replace"              "efficient-programming"
[13] "small-helpers"               "collapse-options"
```

```
R> help("collapse-documentation")
```

This article cannot fully present **collapse**, but the following sections introduce its key features, starting with (2) the *Fast Statistical Functions* and their (3) integration with data manipulation functions; (4) architecture for time series and panel data; (5) table joins and pivots; (6) list processing functions; (7) descriptive tools; and (8) global options. Section 9 provides a small benchmark, Section 10 concludes. For deeper engagement with **collapse**, consult the **documentation resources**, including **vignettes**, **cheatsheet**, **blog**, and **talk with slides**.

³Such as weighted statistics, including several (weighted) quantile and mode estimators, support for fully time-aware computations on irregular series/panels, scaling and centering, advanced (grouped, weighted, panel-decomposed) descriptive statistics etc., all supporting missing values and vectors/matrices/data frames.

⁴As can be seen in the **DuckDB Benchmarks**: **collapse** is highly competitive on the 10-100 million observations datasets, but deteriorates in performance at larger data sizes. There may be performance improvements for very "long data" in the future, but, at present, the treatment of columns as fundamental units of computation (in most cases) is a tradeoff for the highly flexible class-agnostic architecture.

2. Fast statistical functions

The *Fast Statistical Functions*, comprising `fsum()`, `fprod()`, `fmean()`, `fmedian()`, `fmode()`, `fvar()`, `fsd()`, `fmin()`, `fmax()`, `fnth()`, `ffirst()`, `flast()`, `fnobs()`, and `fndistinct()`, are a consistent set of S3-generic statistical functions providing fully vectorized statistical operations in R.⁵ Specifically, operations are vectorized across columns and groups, and may also involve weights or transformations of the input data. Their basic syntax is

```
FUN(x, g = NULL, w = NULL, TRA = NULL, na.rm = TRUE, use.g.names = TRUE, ...)
```

with arguments `x` - data (vector, matrix or data frame-like), `g` - groups (atomic vector, list of vectors, or ‘GRP’ object), `w` - sampling weights (only some functions), and `TRA` - transformation of `x`. The following examples with `fmean()` demonstrate their basic usage on the familiar `iris` dataset recording 50 measurements of 4 variables for 3 species of iris flowers.⁶ All examples support weights (`w`), and `fmean()` can also be multithreaded across columns (`nthreads`).⁷

```
R> fmean(iris$Sepal.Length)
```

```
[1] 5.843
```

```
R> fmean(num_vars(iris))
```

```
Sepal.Length Sepal.Width Petal.Length Petal.Width
      5.843      3.057      3.758      1.199
```

```
R> identical(fmean(num_vars(iris)), fmean(as.matrix(num_vars(iris))))
```

```
[1] TRUE
```

```
R> fmean(iris$Sepal.Length, g = iris$Species)
```

```
setosa versicolor virginica
  5.006      5.936      6.588
```

```
R> fmean(num_vars(iris), g = iris$Species)
```

```
      Sepal.Length Sepal.Width Petal.Length Petal.Width
setosa      5.006      3.428      1.462      0.246
versicolor  5.936      2.770      4.260      1.326
virginica   6.588      2.974      5.552      2.026
```

⁵‘Vectorization’ in R/*collapse* means that these operations are implemented using compiled C/C++ code.

⁶`num_vars()` returns the numeric variables in a data frame-like object; `cat_vars()` the categorical ones.

⁷Not all functions are multithreaded, and parallelism is implemented differently for different functions, as detailed in the respective function documentation. The default use of single instruction multiple data (SIMD) parallelism also implies limited gains from multithreading for simple (non-grouped) operations.

2.1. Transformations

The TRA argument toggles (grouped) replacing and sweeping operations (by reference), generalizing `sweep(x, 2, STATS = fmean(x))`.⁸ Table 1 lists the 11 possible TRA operations.

<i>String</i>	<i>Description</i>
"replace_na"/"na"	replace missing values in <code>x</code> by <code>STATS</code>
"replace_fill"/"fill"	replace data and missing values in <code>x</code> by <code>STATS</code>
"replace"	replace data by <code>STATS</code> but preserve missing values in <code>x</code>
"-"	subtract <code>STATS</code> (center)
"-+"	subtract <code>STATS</code> and add overall average statistic
"/"	divide by <code>STATS</code> (scale)
"%"	compute percentages (divide and multiply by 100)
"+"	add <code>STATS</code>
"*"	multiply by <code>STATS</code>
"%%"	modulus (remainder from division by <code>STATS</code>)
"-%%"	subtract modulus (make data divisible by <code>STATS</code>)

Table 1: Available TRA argument choices.

For example, option `TRA = "fill"` replaces elements with the corresponding statistics:

```
R> fmean(iris$Sepal.Length, g = iris$Species, TRA = "fill")[c(1:5, 51:55)]

[1] 5.006 5.006 5.006 5.006 5.006 5.936 5.936 5.936 5.936 5.936
```

Additionally, a `set` argument can be passed to *Fast Statistical Functions* to toggle transformation by reference. For example `fmean(iris$Sepal.Length, g = iris$Species, TRA = "fill", set = TRUE)` would modify `Sepal.Length` in-place and return it invisibly.

Having grouping and data transformation functionality directly built into generic statistical functions facilitates and speeds up many common operations. It notably avoids the need to convert atomic objects to data frames for grouped aggregations or transformations. The *Fast Statistical Functions* are complemented by a smaller set of *Data Transformation Functions*, including `TRA()`, infix functions such as `%r+%`, `%c+%`, `%+=%` and `setop()` for row/column-wise arithmetic operations (by reference), `dapply()` to apply functions across rows or columns of matrices/data frames, `BY()` for general split-apply-combine computing, and specialized functions such as `fscale()` or `f[hd]within()` for (grouped, weighted) scaling and centering:

```
R> BY(num_vars(iris), g = iris$Species, FUN = mad) |> head(1)

      Sepal.Length Sepal.Width Petal.Length Petal.Width
setosa      0.2965      0.3706      0.1483          0

R> fscale(iris$Sepal.Length, g = iris$Species) |> fsd(g = iris$Species)

      setosa versicolor virginica
      1          1          1
```

⁸The TRA argument internally calls `TRA(x, STATS, FUN = "-", g = NULL, set = FALSE, ...)`.

3. Integration with data manipulation functions

collapse also provides a broad set of **Fast Data Manipulation Functions**, including `fselect()`, `fsubset()`, `fslice()`, `fgroup_by()`, `fsummarise()`, `fmutate()`, `frename()`, `fcount()`, etc. as optimized analogues to base R/**tidyverse** functions. These are integrated with the *Fast Statistical and Transformation Functions* to enable vectorized operations in a familiar data frame oriented and **tidyverse**-like workflow. I illustrate this using the included World Development Dataset (`wlddev`) recording five key indicators for 216 economies and years 1960-2020.

```
R> fndistinct(wlddev)
```

country	iso3c	date	year	decade	region	income	OECD	PCGDP
216	216	61	61	7	7	4	2	9470
LIFEEX	GINI	ODA	POP					
10548	368	7832	12877					

Below, I track changes in Life Expectancy at the country level, computing the range, average change per year, and the correlation with GDP per Capita in log-differences⁹ while also saving the latest measurements.¹⁰ The `collap()` function called at the end to aggregate the country-level statistics across income groups is a convenience function for mixed-type aggregation. By default, it uses `fmean()` for numeric data and `fmode()` for categorical data, here also applying population weights such that the most populous country is selected.¹¹

```
R> wlddev |> fgroup_by(country, income) |>
+   fsummarise(n = fnobs(LIFEEX),
+             diff = fmax(LIFEEX) %-% fmin(LIFEEX),
+             mean_diff = fmean(fdiff(LIFEEX)),
+             cor_PCGDP = pwcor(Dlog(LIFEEX), Dlog(PCGDP)),
+             across(c(LIFEEX, PCGDP, POP), flast)) |>
+   collap( ~ income, w = ~ POP, keep.w = FALSE) |> print(digits = 2)
```

	country	income	n	diff	mean_diff	cor_PCGDP	LIFEEX	PCGDP
1	United States	High income	60	13	0.23	-0.00828	81	44599
2	Ethiopia	Low income	60	26	0.43	0.27153	64	691
3	India	Lower middle income	60	24	0.41	-0.15366	69	2324
4	China	Upper middle income	60	27	0.45	0.00077	76	8749

With the exception of `pwcor(Dlog(LIFEEX), Dlog(PCGDP))`, which is evaluated for every country, all other expressions in `fsummarise()` and also the final `collap()` call are fully vectorized, i.e., *Fast Statistical and Transformation Functions* are evaluated only once with grouping information passed to their `g` arguments. Using subtraction by reference in the range is more memory efficient as the executed expression is equivalent to `fmax(LIFEEX, country) %-% fmin(LIFEEX, country)`. `across()` directly invokes `flast.data.frame()` on the subset of columns. *Fast Statistical Functions* also have a method for grouped data, thus `fsummarise()` is not always needed. Below, I average recent data with population weights.¹²

⁹`Dlog()` abbreviates `fdiff(log(x))`, `pwcor()` wraps `cor(..., use = "pairwise.complete.obs")`.

¹⁰Note that `fsummarise()` never re-groups data, making a call to `fungroup()` redundant.

¹¹See also [this blog post on aggregating survey data using collapse](#) which showcases more aspects of the `collap()` function using real census data.

¹²Functions like `fmean()`, when called on a grouped data frame with weights, by default (`keep.w = TRUE`) sum the weights column and retain it after the grouping columns to permit further weighted operations.

```
R> wlddev |> fsubset(year >= 2015, income, PCGDP:POP) |>
+   fgroup_by(income) |> fmean(POP)

      income    sum.POP PCGDP LIFEEX  GINI      ODA
1      High income 5.901e+09 43340  80.70 36.14  81398948
2      Low income 3.296e+09   663  63.05 39.13 2466732035
3 Lower middle income 1.490e+10  2177  68.31 36.48 2234540565
4 Upper middle income 1.318e+10  8168  75.51 41.68 -86630117
```

Employing the *Fast Statistical Functions* directly often results in more efficient code. For example, computing population shares for each year as `fmutate(wlddev, POP_share = fsum(POP, year, TRA = "/"))` is considerably more efficient than `wlddev |> fgroup_by(year) |> fmutate(POP_share = fsum(POP, TRA = "/")) |> fungroup()`, which in turn is more efficient than using `POP / fsum(POP)` (extra full-length allocation) or `proportions(POP)` (slow split-apply-combine evaluation logic) inside the grouped `fmutate()` call.

collapse thus offers several ways to reach the same outcome. This may be confusing at first, but all options can be distinguished in terms of efficiency. Hence, it is helpful to think about computer resources and attempt to craft minimalistic solutions to complex problems. For example, I recently combined multiple spatial datasets on points of interest (POIs) and needed to deduplicate them. I decided to keep the richest source for each location and POI type/category. After creating comparable POI confidence, location, and type indicators, my deduplication expression ended up being a single line of the form `fsubset(data, source == fmode(source, list(location, type), confidence, "fill"))`—which retains POIs from the confidence-weighted most frequent (richest) source by location and type. This is efficient because it avoids materializing intermediate datasets and relegates all computations to `fmode()`—*Fast Statistical Functions* call a highly optimized internal grouping algorithm.

3.1. Vectorizations for advanced tasks

`fsummarise()` and `fmutate()` follow an eager vectorization approach, such that using any *Fast Statistical Function* in an expression causes the entire expression to be vectorized (i.e., evaluated only once with *Fast Statistical Function* returning grouped output). This only applies to visible expressions, **collapse** cannot read the contents of custom functions, thus such functions are always evaluated using split-apply-combine logic. This eager vectorization approach enables efficient grouped calculations of more complex statistics. Below, I forecast the population of each region via linear regression ($POP \sim year$) in a fully vectorized way.

```
R> wlddev |> collap(POP ~ region + year, FUN = fsum) |>
+   fmutate(POP = POP / 1e6) |> fgroup_by(region) |>
+   fmutate(dmy = fmean(year, TRA = "-")) |>
+   fsummarise(beta = fsum(POP, dmy) %/=% fsum(dmy, dmy),
+             POP20 = flast(POP)) |>
+   fmutate(POP21 = POP20 + beta, POP22 = POP21 + beta,
+           POP23 = POP22 + beta, POP24 = POP23 + beta)

      region  beta POP20 POP21 POP22 POP23 POP24
1 East Asia & Pacific 18.731 2291.4 2310.2 2328.9 2347.6 2366.4
```


2	Europe & Central Asia	2.463	921.2	923.7	926.1	928.6	931.0
3	Latin America & Caribbean	6.460	646.4	652.9	659.4	665.8	672.3
4	Middle East & North Africa	5.409	456.7	462.1	467.5	472.9	478.3
5	North America	2.301	365.9	368.2	370.5	372.8	375.1
6	South Asia	19.640	1835.8	1855.4	1875.1	1894.7	1914.3
7	Sub-Saharan Africa	12.852	1103.5	1116.3	1129.2	1142.0	1154.9

When `fsummarise()` evaluates an expression involving *Fast Statistical Functions*, it sets their `g` argument with a grouping ('GRP') object that is directly handed to C/C++, and also sets `use.g.names = FALSE`. Hence, weights (`w`) becomes the second positional argument. Similarly, `fmutate()` sets `g` and `TRA = "fill"`, which can be overwritten by the user (here with `TRA = "-"`). The expression `fsum(x, dmy) %/=% fsum(dmy, dmy)` amounts to `cov(x, y)/var(y)`, but is vectorized across groups and memory efficient—leveraging the weights (`w`) argument to `fsum()` to compute products (`v * dmy` and `dmy * dmy`) internally and division by reference (`%/=%`) to avoid an additional allocation. In a 2023 [blog post](#), I forecasted high-resolution population estimates for South Africa like this. Using 1 km^2 [WorldPop](#) data available for years 2014-2020, I ran 1.6 million cell-regressions and obtained 2 forecasts for 2021 and 2022 in less than 0.3 seconds on my M1 Mac. Another neat example from the community, shared by Andrew Ghazi in a [blog post](#), vectorizes an expression to compute the p value, `2 * pt(abs(fmean(x) * sqrt(6) / fsd(x)), 5, lower.tail = FALSE)`, across 300,000 groups for a simulation study, yielding a 70x performance increase over `dplyr`.

collapse also vectorizes advanced statistics. The following calculates a weighted set of summary statistics by groups, with weighted quantiles type 8 following [Hyndman and Fan \(1996\)](#).¹³

```
R> wlddev |> fsubset(is.finite(POP)) |> fgroup_by(region) |>
+   fmutate(o = radixorder(GRPid(), LIFEEX)) |>
+   fsummarise(min = fmin(LIFEEX),
+             Q1 = fnth(LIFEEX, 0.25, POP, o = o, ties = "q8"),
+             mean = fmean(LIFEEX, POP),
+             median = fmedian(LIFEEX, POP, o = o),
+             Q3 = fnth(LIFEEX, 0.75, POP, o = o, ties = "q8"),
+             max = fmax(LIFEEX))
```

	region	min	Q1	mean	median	Q3	max
1	East Asia & Pacific	18.91	65.28	68.45	69.67	73.86	85.08
2	Europe & Central Asia	45.37	68.68	72.30	71.58	76.67	85.42
3	Latin America & Caribbean	41.76	65.17	69.16	70.87	74.48	82.19
4	Middle East & North Africa	29.92	61.96	66.65	69.12	72.64	82.80
5	North America	68.90	73.57	75.54	75.62	78.38	82.05
6	South Asia	32.45	55.08	60.19	62.00	66.67	78.92
7	Sub-Saharan Africa	26.17	46.51	52.53	52.23	58.32	74.51

Weighted quantiles have a sub-column level parallel implementation,¹⁴ and, as shown above, can also harness an (optional) optimization via an overall ordering vector—combining groups with the data column to avoid repeated sorting of the same elements by different functions.

¹³**collapse** computes weighted quantiles in a theoretically consistent way, see [fquantile](#) for details.

¹⁴Use `set_collapse(nthreads = #)` or set the `nthreads` arguments of `fnth()`/`fmedian()` (default 1).

3.2. Grouping objects and lower-level API

Whereas the `g` argument supports ad-hoc grouping with vectors and lists/data frames, for repeated operations the cost of grouping can be minimized by using factors (see `?qF` for efficient factor generation) or ‘GRP’ objects as inputs. The latter contain all information **collapse**’s statistical functions may require to operate across groups and are thus passed to internal C/C++ code without checks. They can be created with `GRP()`. Its basic syntax is:

```
GRP(X, by = NULL, sort = TRUE, return.groups = TRUE, method = "auto", ...)
```

Below, I create a ‘GRP’ object from two columns in the World Development Dataset (`wlddev`). The `by` argument also supports column names/indices, and `X` could also be an atomic vector.

```
R> str(g <- GRP(wlddev, ~ income + OECD))
```

```
Class 'GRP'  hidden list of 9
 $ N.groups      : int 6
 $ group.id      : int [1:13176] 3 3 3 3 3 3 3 3 3 3 ...
 $ group.sizes   : int [1:6] 2745 2074 1830 2867 3538 122
 $ groups        : 'data.frame':      6 obs. of  2 variables:
 ..$ income: Factor w/ 4 levels "High income",...: 1 1 2 3 4 4
 ..$ OECD   : logi [1:6] FALSE TRUE FALSE FALSE FALSE TRUE
 ..$ attr(*, "label")= chr "Income Level"
 ..$ attr(*, "label")= chr "Is OECD Member Country?"
 $ group.vars    : chr [1:2] "income" "OECD"
 $ ordered       : Named logi [1:2] TRUE FALSE
 ..$ attr(*, "names")= chr [1:2] "ordered" "sorted"
 $ order         : int [1:13176] 245 246 247 248 249 250 251 252 253 254 ...
 ..$ attr(*, "starts")= int [1:6] 1 2746 4820 6650 9517 13055
 ..$ attr(*, "maxgrp")= int 3538
 ..$ attr(*, "sorted")= logi FALSE
 $ group.starts  : int [1:6] 245 611 1 306 62 7687
 $ call          : language GRP.default(X = wlddev, by = ~income + OECD)
```

‘GRP’ objects make grouped statistical computations in **collapse** fully programmable. I can employ the object with the *Fast Statistical Functions* and some utilities¹⁵ to efficiently aggregate GDP per capita, life expectancy, and country name, again applying population weights.

```
R> add_vars(g$groups,
+   get_vars(wlddev, "country") |> fmode(g, wlddev$POP, use = FALSE),
+   get_vars(wlddev, c("PCGDP", "LIFEEX")) |> fmean(g, wlddev$POP, use = F),
+   get_vars(wlddev, "POP") |> fsum(g, use = FALSE))
```

	income	OECD	country	PCGDP	LIFEEX	POP
1	High income	FALSE	Saudi Arabia	22426.7	73.00	3.114e+09

¹⁵`add_vars()` is a fast `cbind.data.frame()` which also has an assignment method, and `get_vars()` enables fast and secure extraction of data frame columns. `use = FALSE` abbreviates `use.g.names = FALSE`.

```

2      High income  TRUE United States 31749.6  75.84 5.573e+10
3      Low income  FALSE      Ethiopia   557.1  53.51 2.095e+10
4 Lower middle income FALSE      India  1238.8  60.59 1.138e+11
5 Upper middle income FALSE      China  3820.6  68.21 1.114e+11
6 Upper middle income  TRUE      Mexico  8311.2  69.06 8.162e+09

```

The above is equivalent to `collap(wlddev, country + PCGDP + LIFEEX ~ income + OECD, w = ~ POP)`, which internally toggles many of the same function calls.

Similarly, data can be transformed, here using `fwithin()` to level average differences in economic status, adding back the overall mean after subtracting out group means:¹⁶

```

R> add_vars(wlddev) <- get_vars(wlddev, c("PCGDP", "LIFEEX")) |>
+   fwithin(g, mean = "overall.mean") |> add_stub("center_")

```

The lower-level API is useful for package development and standard-evaluation programming, as further elucidated in the [vignette on developing with collapse](#). Users should note that **collapse** does not provide metaprogramming capabilities in its non-standard evaluation functions—such as [quosures](#) or [indirection](#) familiar to **tidyverse** users. Instead, it has standard-evaluation equivalents to some of these functions which typically end with a `v` for ‘variables’, such as `collapv()`, `fslicev()`, `ftransformv()`, etc. In other cases, including those typically handled by `fsummarise()` or `fmutate()`, users need to use the lower-level API for programming, or resort to `substitute()` and friends. The main reason for not providing higher-level metaprogramming capabilities is to keep all functions as simple as possible. It also compels users to think deeply about their programs and devise more efficient solutions.

As the small benchmark below illustrates, **dplyr**’s internally more complex data manipulation functions produce much greater overheads, which can noticeably add-up in longer scripts.¹⁷

```

R> bmark(collapse = collapse::fsummarise(wlddev, mean = fmean(PCGDP)),
+       dplyr = dplyr::summarise(wlddev, mean = fmean(PCGDP)))

```

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	collapse	6.68µs	8.04µs	1.05KB	9999	1	83.59ms
2	dplyr	215.13µs	254.04µs	1.54MB	9936	64	2.78s

Grouped programming using ‘GRP’ objects and *Fast Statistical Functions* is also particularly powerful with vectors and matrices. For example, in the [useR 2022 presentation](#), I aggregate 32 global input-output tables stored as matrices (`x`) from the country to the region level using a single grouping object (`g`) and expressions like `x |> fsum(g) |> t() |> fsum(g) |> t()`—computing 45 million sums crunching 5.7GB of data in ~0.3 seconds on my M1.¹⁸

¹⁶`add_stub()` adds a prefix (or suffix if `pre = FALSE`) to columns \rightarrow `center_PCGDP` and `center_LIFEEX`.

¹⁷`bmark()` is a slim wrapper around `bench::mark()`. See the Computational details section.

¹⁸Another recent application with vectors involved numerically optimizing a parameter a in an equation of the form $\sum_i x_{ij}^a \forall j \in J$ so as to minimize the deviation from a target y_j where there are J groups (1 million in my case)—see the first example in [this blog post](#) for an illustration.

4. Time series and panel series

collapse also provides a flexible high-performance architecture to perform time aware computations on time series and panel series. In particular, users can either apply time series and panel data transformations by passing individual and/or time identifiers to the respective functions in an ad-hoc fashion, or by using ‘**indexed_frame**’ and ‘**indexes_series**’ classes, which implement full and deep indexation for worry-free application in many contexts. Table 2 compactly summarizes **collapse**’s architecture for time series and panel data.

Classes, constructors and utilities

`findindex_by()`, `findindex()`, `unindex()`, `reindex()`, `timeid()`, `is_irregular()`,
`to_plm()` + S3 methods for ‘**indexed_frame**’, ‘**indexed_series**’ and ‘**index_df**’

Core time-based functions

`flag()`, `fdiff()`, `fgrowth()`, `fcumsum()`, `psmat()`
`psacf()`, `pspacf()`, `psccf()`

Data transformation functions with supporting methods

`fscale()`, `f[hd]between()`, `f[hd]within()`

Data manipulation functions with supporting methods

`fsubset()`, `funique()`, `roworder[v]()` (internal), `na_omit()` (internal)

Summary functions with supporting methods

`varying()`, `qsu()`

Table 2: Time series and panel data architecture.

4.1. Ad-hoc computations

Time series functions such as `fgrowth()` to compute growth rates are S3 generic and can be applied to most time series classes. In addition to a `g` argument for grouping, these functions also have a `t` argument for indexation. But first, I provide a basic example of computing the annualized 10-year growth rates in miles flown by airlines in the USA from 1937-1960.

```
R> fgrowth(airmiles, n = 10, power = 1/10) |> na.omit() |> round(1)
```

Time Series:

Start = 1947

End = 1960

Frequency = 1

```
[1] 31.0 28.7 25.7 22.5 22.5 24.3 24.6 22.6 19.4 14.2 15.3 15.5 15.8 14.3
```

The results show that the flight volume has been growing steadily but at a decreasing rate.

To illustrate the full capabilities of these time series functions, I generate a sector-level trade dataset of export values (`v`) by country (`c`), sector (`s`), and year (`y`). Like many detailed trade datasets, it is unbalanced—not all sectors/products are exported by country `c` in all years.

```
R> set.seed(101)
R> exports <- expand.grid(y = 2001:2010, c = paste0("c", 1:10),
+   s = paste0("s", 1:10)) |> fmutate(v = abs(rnorm(1e3))) |>
+   colorder(c, s) |> fsubset(-sample.int(1e3, 500))
```

The following extracts one country-sector series from the `exports` dataset. It is irregular, missing years 2003 and 2006.¹⁹ Indexation using the `t` argument still allows for correct (time-aware) computations in this context without the need to 'expand' the data/fill gaps.

```
R> .c(y, v) %=% fsubset(exports, c == "c1" & s == "s7", -c, -s)
R> print(y)
```

```
[1] 2001 2002 2004 2005 2007 2008 2009 2010
```

```
R> fgrowth(v, t = y) |> round(2)
```

```
[1]      NA 175.52      NA -22.37      NA 624.27 -79.01 534.56
```

```
R> fgrowth(v, -1:3, t = y) |> head(4)
```

```
      FG1      --      G1    L2G1    L3G1
[1,] -63.71 0.3893      NA      NA      NA
[2,]      NA 1.0726 175.52      NA      NA
[3,]  28.82 0.8450      NA -21.22 117.05
[4,]      NA 0.6559 -22.37      NA -38.85
```

If `t` is a plain numeric vector or factor as in this case, it is coerced to integer and interpreted as time steps.²⁰ Time objects like 'Date' or 'POSIXct' on the other hand are internally passed through `timeid()` to generate an appropriate integer representation of them.²¹

Functions `flag()/fdiff()/fgrowth()` also have associated 'operators' `L()/D()/G()` to facilitate their use inside formulas and provide an enhanced data frame interface for convenient ad-hoc computations. With panel data, `t` can be omitted, but this requires sorted data with consecutive groups. Below, I demonstrate two ways to compute a sequence of lagged growth rates using either `G()` or `fgrowth()` and `tfrm()`—a shorthand for `ftransform()`.²²

```
R> G(exports, -1:2, by = v ~ c + s, t = ~ y) |> head(3)
```

```
      c  s    y  FG1.v      v  G1.v L2G1.v
1 c1 s1 2002 -18.15 0.5525      NA      NA
2 c1 s1 2003 214.87 0.6749 22.17      NA
3 c1 s1 2004 -31.02 0.2144 -68.24 -61.2
```

¹⁹`%=%` is an infix operator for the `massign()` function in **collapse** which is a multivariate version of `assign()`.

²⁰This is premised on the observation that the most common form of temporal identifier is a plain numeric vector representing calendar years. Users should manually call `timeid()` on plain vectors with decimals.

²¹`timeid()` divides by the greatest common divisor (GCD) and subtracts the minimum to generate an integer-id (starting from 1). For this approach to work, `t` must have an appropriate class, e.g., for monthly/quarterly data, `zoo::yearmon()/zoo::yearqtr()` should be used instead of 'Date' or 'POSIXct'.

²²Several key functions in **collapse** have syntactic shorthands. The `list(v = v)` is needed here to prevent `fgrowth()` from creating a matrix with the growth rates, ensuring that the 'list' method applies.

```
R> tfm(exports, fgrowth(list(v = v), -1:2, g = list(c, s), t = y)) |> head(3)
```

```
   c  s    y      v FG1.v  G1.v L2G1.v
1 c1 s1 2002 0.5525 -18.15    NA     NA
2 c1 s1 2003 0.6749 214.87  22.17    NA
3 c1 s1 2004 0.2144 -31.02 -68.24 -61.2
```

These functions and operators are also integrated with `fgroup_by()` and `fmutate()` for vectorized computations. As mentioned earlier, ad-hoc grouping is always more efficient.

```
R> A <- exports |> fgroup_by(c, s) |> fmutate(gv = G(v, t = y)) |> fungroup()
R> head(B <- exports |> fmutate(gv = G(v, g = list(c, s), t = y)), 4)
```

```
   c  s    y      v    gv
1 c1 s1 2002 0.5525    NA
2 c1 s1 2003 0.6749  22.17
3 c1 s1 2004 0.2144 -68.24
4 c1 s1 2005 0.3108  44.98
```

```
R> identical(A, B)
```

```
[1] TRUE
```

4.2. Indexed series and frames

For more complex use cases, indexation is convenient. **collapse** supports **plm**'s `'pseries'` and `'pdata.frame'` classes through dedicated methods. Flexibility and performance considerations lead to the creation of new classes `'indexes_series'` and `'indexed_frame'` which inherit from the former. Any data frame-like object can become an `'indexed_frame'` and behave as usual for other operations. The technical implementation of these classes is described in the [vignette on object handling](#) and, in more detail, in the [documentation](#). Their basic usage is:

```
data_ix <- findex_by(data, id1, ..., time)
data_ix$indexed_series; with(data, indexed_series)
index_df <- findex(data_ix)
```

Data can be indexed using one or more indexing variables. Unlike `'pdata.frame'`, an `'indexed_frame'` is a deeply indexed structure—every series inside the frame is already an `'indexes_series'`. A comprehensive set of [methods for subsetting and manipulation](#), and applicable `'pseries'` and `'pdata.frame'` methods for time series and transformation functions like `flag()/L()` ensure that these objects behave in a time-/panel-aware manner in any caller environment (`with()`, `lm()`, etc.). Indexation can be undone using `unindex()` and redone with `reindex()` and a suitable `'index_df'`. `'indexes_series'` can be atomic vectors or matrices (including `'ts'` or `'xts'`) and can also be created directly using `reindex()`.

```
data <- unindex(data_ix)
data_ix <- reindex(data, index = index_df)
indexed_series <- reindex(vec/mat, index = vec/index_df)
```

It is worth highlighting that the flexibility of this architecture is new to the R ecosystem: A ‘pdata.frame’ or ‘fixest_panel’ only works inside **plm/fixest** estimation functions.²³ Time series classes like ‘xts’ and ‘tsibble’ also do not provide deeply indexed structures or native handling of irregularity in basic operations. ‘indexed_series’ and ‘indexed_frame’, on the other hand, work ‘anywhere’, and can be superimposed on any suitable object as long as **collapse**’s functions (`flag()/L()`, etc.) are used to perform the time-based computations.

An example follows using the `exports` data. *Note* that data can be unsorted for indexation.

```
R> exportsi <- exports /> findex_by(c, s, y)
R> exportsi /> G(0:1) /> head(5)
```

```
      c  s    y      v   G1.v
1 c1 s1 2002 0.5525    NA
2 c1 s1 2003 0.6749  22.17
3 c1 s1 2004 0.2144 -68.24
4 c1 s1 2005 0.3108  44.98
5 c1 s1 2006 1.1740 277.76
```

```
Indexed by:  c.s [1] | y [5 (10)]
```

```
R> exportsi /> findex() /> print(2)
```

```
      c.s    y
1 c1.s1 2002
2 c1.s1 2003
---
499 c10.s10 2007
500 c10.s10 2009
```

```
c.s [100] | y [10]
```

The index statistics are: [N. ids] | [N. periods (total periods: (max-min)/GCD)].

```
R> vi <- exportsi$v; str(vi, width = 70, strict = "cut")
```

```
'indexed_series' num [1:500] 0.552 0.675 0.214 0.311 1.174 ...
- attr(*, "index_df")=Classes 'index_df', 'pindex' and 'data.frame'..
..$ c.s: Factor w/ 100 levels "c1.s1","c2.s1",...: 1 1 1 1 1 1 1 1 ..
..$ y : Ord.factor w/ 10 levels "2001"<"2002"<...: 2 3 4 5 6 7 8 9..
..- attr(*, "nam")= chr [1:3] "c" "s" "y"
```

```
R> is_irregular(vi)
```

```
[1] TRUE
```

²³And, in the case of **fixest**, inside **data.table** due to dedicated methods.

```
R> vi /> psmat() /> head(3)
```

	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010
c1.s1	NA	0.552	0.675	0.214	0.311	1.17	0.619	0.1127	0.917	0.223
c2.s1	NA	0.795	NA	NA	0.237	NA	NA	0.0585	0.818	NA
c3.s1	NA	0.709	0.268	1.464	NA	NA	0.467	0.1193	0.467	NA

```
R> fdiff(vi) /> psmat() /> head(3)
```

	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010
c1.s1	NA	NA	0.122	-0.461	0.0964	0.863	-0.555	-0.506	0.804	-0.694
c2.s1	NA	NA	NA	NA	NA	NA	NA	NA	0.759	NA
c3.s1	NA	NA	-0.441	1.196	NA	NA	NA	-0.348	0.348	NA

`psmat()`, for panel-series to matrix, generates a matrix or array from panel data. Thanks to deep indexation, indexed computations work inside arbitrary data masking environments:

```
R> settransform(exportsi, v_ld = Dlog(v))
R> lm(v_ld ~ L(v_ld, 1:2), exportsi) /> summary() /> coef() /> round(3)
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.008	0.141	-0.058	0.954
L(v_ld, 1:2)L1	-0.349	0.115	-3.042	0.004
L(v_ld, 1:2)L2	-0.033	0.154	-0.215	0.831

Indexed series/frames also support transformations such as grouped scaling with `fscale()` or demeaning with `fwthin()`. Functions `psacf()`/`pspacf()`/`psccf()` provide panel-data autocorrelation functions, which are computed using group-scaled and suitably lagged panel-series. The ‘`index_df`’ attached to these objects can also be used with other general tools such as `collapse::BY()` to perform grouped computations using third-party functions:

```
R> BY(vi, findex(vi)$c.s, data.table::frollmean, 5) /> head(10)
```

[1]	NA	NA	NA	NA	0.5853	0.5986	0.4861	0.6267	0.6092	NA
-----	----	----	----	----	--------	--------	--------	--------	--------	----

```
Indexed by: c.s [2] | y [9 (10)]
```

Last but not least, the computational performance of these classes is second to none.²⁴

5. Table joins and pivots

Among its suite of **data manipulation functions**, **collapse**’s implementations of table **joins** and **pivots** are particularly noteworthy since they offer several new features, including rich verbosity for table joins, pivots supporting variable labels, and recast pivots. Both implementations also offer outstanding computational performance, syntax, and memory efficiency.

²⁴See, e.g., the small benchmark presented on slide 40 of the [useR 2022 presentation](#).

5.1. Joins

Compared to commercial software such as STATA (StataCorp LLC. 2023), the implementation of joins in most open-source software provides no information on how many records were joined from both tables. This often provokes manual efforts to validate the join operation. `collapse::join()` provides many options to understand table join operations. Its syntax is:

```
join(x, y, on = NULL, how = "left", suffix = NULL, validate = "m:m",
     multiple = FALSE, sort = FALSE, keep.col.order = TRUE, verbose = 1,
     drop.dup.cols = FALSE, require = NULL, column = NULL, attr = NULL, ...)
```

It defaults to left join and only takes first matches from `y` (`multiple = FALSE`). Thus, it simply adds columns to `x`, which is efficient and sufficient/desired in many cases. By default (`verbose = 1`), it prints information about the join operation and number of records joined.

To demonstrate `join()`, I generate a small database for a bachelor in economics curriculum. It has a `teacher` table of 4 teachers (`id`: PK) and a linked (`id`: FK) `course` table of 5 courses.

```
R> teacher <- data.frame(id = 1:4, names = c("John", "Jane", "Bob", "Carl"),
+   age = c(35, 32, 42, 67), subject = c("Math", "Econ", "Stats", "Trade"))
R> course <- data.frame(id = c(1, 2, 2, 3, 5), semester = c(1, 1, 2, 1, 2),
+   course = c("Math I", "Microecon", "Macroecon", "Stats I", "History"))
R> join(teacher, course, on = "id")
```

```
left join: teacher[id] 3/4 (75%) <1:1st> course[id] 3/5 (60%)
  id names age subject semester  course
1  1  John  35   Math         1   Math I
2  2  Jane  32   Econ         1 Microecon
3  3   Bob  42  Stats         1   Stats I
4  4  Carl  67  Trade         NA    <NA>
```

Users can request the generation of a `.join` column (`column = "name"/TRUE`) akin to STATA's `_merge` column to indicate the origin of records in the joined table—useful with a full join:

```
R> join(teacher, course, how = "full", multiple = TRUE, column = TRUE)
```

```
full join: teacher[id] 3/4 (75%) <1:1.33> course[id] 4/5 (80%)
  id names age subject semester  course  .join
1  1  John  35   Math         1   Math I matched
2  2  Jane  32   Econ         1 Microecon matched
3  2  Jane  32   Econ         2 Macroecon matched
4  3   Bob  42  Stats         1   Stats I matched
5  4  Carl  67  Trade         NA    <NA> teacher
6  5 <NA>  NA   <NA>         2   History  course
```

An alternative is to request an attribute (`attr = "name"/TRUE`) that also summarizes the join operation, including the output of `fmatch()`—the workhorse of `join()` if `sort = FALSE`.

```
R> join(teacher, course, multiple = TRUE, attr = "jn") |> attr("jn") |> str()

left join: teacher[id] 3/4 (75%) <1:1.33> course[id] 4/5 (80%)
List of 3
 $ call      : language join(x = teacher, y = course, multiple = TRUE, "..
 $ on.cols:List of 2
  ..$ x: chr "id"
  ..$ y: chr "id"
 $ match     : 'qG' int [1:5] 1 2 3 4 NA
  ..- attr(*, "N.nomatch")= int 1
  ..- attr(*, "N.groups")= int 5
  ..- attr(*, "N.distinct")= int 4
```

Users can also invoke the `validate` argument to examine the uniqueness of the join keys in either table: passing a '1' produces an error if the respective key is not unique.

```
R> join(teacher, course, on = "id", validate = "1:1") |>
+   tryCatch(error = function(e) strwrap(e) |> cat(sep = "\n"))
```

```
Error in join(teacher, course, on = "id", validate = "1:1"): Join is
not 1:1: teacher (x) is unique on the join columns; course (y) is
not unique on the join columns
```

Similarly, the `require` argument allows users to demand a minimum matching success rate.

```
R> join(teacher, course, on = "id", require = list(x = 0.8)) |>
+   tryCatch(error = function(e) substr(e, 102, 200) |> cat())
```

```
Matched 75.0% of records in table teacher (x), but 80.0% is required
```

A few further particularities are worth highlighting. First, `join()` is class-agnostic and preserves the attributes of `x` (any list-based object). It supports 6 different join operations ("left", "right", "inner", "full", "semi", or "anti"). This demonstrates the latter two:

```
R> for (h in c("semi", "anti")) join(teacher, course, how = h) |> print()
```

```
semi join: teacher[id] 3/4 (75%) <1:1st> course[id] 3/5 (60%)
  id names age subject
1  1  John  35    Math
2  2  Jane  32    Econ
3  3   Bob  42    Stats
anti join: teacher[id] 3/4 (75%) <1:1st> course[id] 3/5 (60%)
  id names age subject
1  4  Carl  67    Trade
```

By default (`sort = FALSE`), the order of rows in `x` is preserved. Setting `sort = TRUE` sorts all records in the joined table by the keys.²⁵ The join relationship is indicated inside the `<>` as the number of records joined from each table divided by the number of unique matches.

²⁵This is done using a separate sort-merge-join algorithm, so it is faster than performing a hash join using `fmatch()` followed by sorting, particularly if the data is already sorted on the keys.

`join()`'s handling of duplicate columns in both tables is also rather special. By default (`suffix = NULL`), `join()` extracts the name of the `y` table and appends `y`-columns with it.

```
R> course$names <- teacher$names[course$id]
R> join(teacher, course, on = "id", how = "inner", multiple = TRUE)
```

```
inner join: teacher[id] 3/4 (75%) <1:1.33> course[id] 4/5 (80%)
duplicate columns: names => renamed using suffix '_course' for y
```

	id	names	age	subject	semester	course	names_course
1	1	John	35	Math	1	Math I	John
2	2	Jane	32	Econ	1	Microecon	Jane
3	2	Jane	32	Econ	2	Macroecon	Jane
4	3	Bob	42	Stats	1	Stats I	Bob

This is congruent to the principle of adding columns to `x` in the default first-match left join by altering this table as little as possible. Alternatively, option `drop.dup.cols = "x"/"y"` can be used to remove duplicate columns from either `x` or `y` before the join operation.

```
R> join(teacher, course, on = "id", multiple = TRUE, drop.dup.cols = "y")
```

```
left join: teacher[id] 3/4 (75%) <1:1.33> course[id] 4/5 (80%)
duplicate columns: names => dropped from y
```

	id	names	age	subject	semester	course
1	1	John	35	Math	1	Math I
2	2	Jane	32	Econ	1	Microecon
3	2	Jane	32	Econ	2	Macroecon
4	3	Bob	42	Stats	1	Stats I
5	4	Carl	67	Trade	NA	<NA>

A final noteworthy feature is that `fmatch()` has a built-in overidentification check, which warns if more key columns than necessary to identify the records are provided. This check only triggers with 3+ id columns as for efficiency reasons the first two ids are jointly hashed. `join()` is thus a highly efficient, versatile, and verbose implementation of table joins for R.

5.2. Pivots

The reshaping/pivoting functionality of both commercial and open-source software is also presently unsatisfactory for complex datasets such as surveys or disaggregated production, trade, or financial sector data, where variable names resemble codes and variable labels are essential to making sense of the data. Such datasets can presently only be reshaped by losing these labels or manual efforts to retain them. Modern R packages also offer different reshaping functions, such as `data.table::melt()/tidyr::pivot_longer()` to combine columns and `data.table::dcast()/tidyr::pivot_wider()` to expand them, requiring users to learn both. Since the depreciation of `reshape(2)` (Wickham 2007), there is also no modern replacement for `reshape2::recast()`, requiring R users to consecutively call two functions.

`collapse::pivot()` provides a powerful new implementation of reshaping for R addressing these shortcomings. It has a single intuitive syntax to perform 'longer', 'wider', and 'recast' pivots and supports complex labelled data without loss of information. Its basic syntax is:

```
pivot(data, ids = NULL, values = NULL, names = NULL, labels = NULL,
      how = "longer", na.rm = FALSE, check.dups = FALSE, FUN = "last", ...)
```

The demonstration below employs the included Groningen Growth and Development Centre 10-Sector Database ([GGDC10S](#)) providing long-run internationally comparable data on sectoral productivity performance in Africa, Asia, and Latin America. While the database covers 10 sectors, for the demonstration I only retain Agriculture, Mining, and Manufacturing.²⁶

```
R> data <- GGDC10S />
+   fmutate(Label = ifelse(Variable == "VA", "Value Added", "Employment")) />
+   fsubset(is.finite(AGR), Country, Variable, Label, Year, AGR:MAN)
R> namlab(data, N = TRUE, Ndistinct = TRUE, class = TRUE)
```

	Variable	Class	N	Ndist	Label
1	Country	character	4364	43	Country
2	Variable	character	4364	2	Variable
3	Label	character	4364	2	<NA>
4	Year	numeric	4364	67	Year
5	AGR	numeric	4364	4353	Agriculture
6	MIN	numeric	4355	4224	Mining
7	MAN	numeric	4355	4353	Manufacturing

To reshape this dataset into a longer format, it suffices to call `pivot(data, ids = 1:4)`. If `labels = "name"` is specified, variable labels stored in `attr(column, "label")` are saved to an additional column. In addition, `names = list(variable = "var_name", value = "val_name")` can be passed to set alternative names for the variable and value columns.

```
R> head(dl <- pivot(data, ids = 1:4, names = list("Sectorcode", "Value"),
+   labels = "Sector", how = "longer"))
```

	Country	Variable	Label	Year	Sectorcode	Sector	Value
1	BWA	VA	Value Added	1964	AGR	Agriculture	16.30
2	BWA	VA	Value Added	1965	AGR	Agriculture	15.73
3	BWA	VA	Value Added	1966	AGR	Agriculture	17.68
4	BWA	VA	Value Added	1967	AGR	Agriculture	19.15
5	BWA	VA	Value Added	1968	AGR	Agriculture	21.10
6	BWA	VA	Value Added	1969	AGR	Agriculture	21.86

`pivot()` only requires essential information and intelligently guesses the rest. For example, the same result could have been obtained by specifying `values = c("AGR", "MIN", "MAN")` instead of `ids = 1:4`. An exact reverse operation can also be specified as `pivot(dl, 1:4, "Value", "Sectorcode", "Sector", how = "wider")`, where `dl` is the long data.

The second option is a wider pivot with `how = "wider"`. Here, `names` and `labels` can be used to select columns containing the names of new columns and their labels.²⁷ Note below

²⁶The "Label" column is added for demonstration purposes. `namlab()` provides a compact overview of variable names and labels stored in `attr(column, "label")`, with (optional) additional information/statistics.

²⁷If multiple columns are selected, they are combined using "_" for names and "-" for labels.

how the labels are combined with existing labels such that also this operation is without loss of information. It is, however, a destructive operation, as with two or more columns selected through `values`, `pivot()` is not able to reverse it. Further arguments like `na.rm`, `fill`, `drop`, `sort`, and `transpose` can be invoked to control the casting process/output.

```
R> head(dw <- pivot(data, c("Country", "Year"), names = "Variable",
+   labels = "Label", how = "wider"))
```

	Country	Year	AGR_VA	AGR_EMP	MIN_VA	MIN_EMP	MAN_VA	MAN_EMP
1	BWA	1964	16.30	152.1	3.494	1.9400	0.7366	2.420
2	BWA	1965	15.73	153.3	2.496	1.3263	1.0182	2.330
3	BWA	1966	17.68	153.9	1.970	1.0022	0.8038	1.282
4	BWA	1967	19.15	155.1	2.299	1.1192	0.9378	1.042
5	BWA	1968	21.10	156.2	1.839	0.7855	0.7503	1.069
6	BWA	1969	21.86	157.4	5.245	2.0314	2.1396	2.124

```
R> namlab(dw)
```

	Variable	Label
1	Country	Country
2	Year	Year
3	AGR_VA	Agriculture - Value Added
4	AGR_EMP	Agriculture - Employment
5	MIN_VA	Mining - Value Added
6	MIN_EMP	Mining - Employment
7	MAN_VA	Manufacturing - Value Added
8	MAN_EMP	Manufacturing - Employment

For the recast pivot (`how = "recast"`), unless a column named `variable` exists in the data, the source and (optionally) destination of variable names needs to be specified using a list passed to `names`, and similarly for `labels`. Again, taking along labels is entirely optional—omitting either the labels-list's `from` or `to` elements will omit the respective operation.

```
R> head(dr <- pivot(data, c("Country", "Year"),
+   names = list(from = "Variable", to = "Sectorcode"),
+   labels = list(from = "Label", to = "Sector"), how = "recast"))
```

	Country	Year	Sectorcode	Sector	VA	EMP
1	BWA	1964	AGR	Agriculture	16.30	152.1
2	BWA	1965	AGR	Agriculture	15.73	153.3
3	BWA	1966	AGR	Agriculture	17.68	153.9
4	BWA	1967	AGR	Agriculture	19.15	155.1
5	BWA	1968	AGR	Agriculture	21.10	156.2
6	BWA	1969	AGR	Agriculture	21.86	157.4

```
R> vlabels(dr)[3:6]
```

Sectorcode	Sector	VA	EMP
NA	NA	"Value Added"	"Employment"

This (`dr`) is the tidy format (Wickham 2014) where each variable is a separate column. It is analytically more useful, e.g., to compute labor productivity as `settransform(dr, LP = VA / EMP)` or to estimate a panel-regression with sector fixed-effects. The recast pivot is thus a natural operation to change data representation. As with the other pivots, it preserves all information and can be reversed by simply swapping the contents of the `from` and `to` keywords.

`pivot()` also supports fast aggregation pivots, the default being `FUN = "last"`, which simply overwrites values in appearance order if the combination of `ids` and `names` does not fully identify the data. The latter can be checked with `check.dups = TRUE`. A small number extremely fast internal aggregation functions: `"first"`, `"last"`, `"sum"`, `"mean"`, `"min"`, `"max"`, and `"count"`, operate 'on the fly' during reshaping. `pivot()` also supports *Fast Statistical Functions*, which will yield vectorized aggregations, but requires a deep copy of the columns aggregated which is avoided by the internal functions. The following example performs aggregation across years with the internal mean function during a recast pivot.

```
R> head(dr_agg <- pivot(data, "Country", c("AGR", "MIN", "MAN"),
+   how = "recast", names = list(from = "Variable", to = "Sectorcode"),
+   labels = list(from = "Label", to = "Sector"), FUN = "mean"))
```

	Country	Sectorcode	Sector	VA	EMP
1	BWA	AGR	Agriculture	462.2	188.06
2	ETH	AGR	Agriculture	34389.9	17624.34
3	GHA	AGR	Agriculture	1549.4	3016.04
4	KEN	AGR	Agriculture	139705.9	5348.91
5	MWI	AGR	Agriculture	28512.6	2762.62
6	MUS	AGR	Agriculture	3819.6	59.34

The [documentation examples](#) demonstrate more features of `pivot()`. Notably, it can also perform longer and recast pivots without id variables, similar to `data.table::transpose()`.

6. List processing

Often in programming, nested structures are needed. A typical use case involves running statistical procedures for multiple configurations of variables and parameters and saving multiple objects, such as a model predictions and performance statistics, in a list. Nested data is also often the result of web scraping or web APIs. A typical use case in development involves serving different data according to user choices. Except for certain recursive functions in packages such as **purrr**, **tidyr**, or **rrapply**, R lacks a general recursive toolkit to create, query, and tidy nested data. **collapse**'s [list processing functions](#) attempt to provide a basic toolkit.

To create nested data, `rsplit()` generalizes `split()` and (recursively) splits up data frame-like objects into (nested) lists. For example, we can split the `GGDC10S` data by country and variable, such that agricultural employment in Argentina can be accessed as:²⁸

²⁸If a nested structure is not needed, `flatten = TRUE` lets `rsplit()` operate like a faster version of `split()`.

```
R> d_list <- GGDC10S |> rsplit( ~ Country + Variable)
R> d_list$ARG$EMP$AGR[1:12]
```

```
[1] 1800 1835 1731 2030 1889 1843 1789 1724 1678 1725 1650 1553
```

This is a convenient data representation for *Shiny Apps* where we can let the user choose data (e.g., `d_list[[input$country]][[input$variable]][[input$sector]]`) without expensive subsetting operations. As mentioned, such data representation can also be the result of an API call parsing JSON or a nested loop or `lapply()` call. Below, I write a nested loop running a regression of agriculture on mining and manufacturing output and employment.

```
R> results <- list()
R> for (country in c("ARG", "BRA", "CHL")) {
+   for (variable in c("EMP", "VA")) {
+     m <- lm(log(AGR+1) ~ log(MIN+1) + log(MAN+1) + Year,
+           data = d_list[[country]][[variable]])
+     results[[country]][[variable]] <- list(model = m, BIC = BIC(m),
+           summary = summary(m))
+   }
+ }
```

This programming may not be ideal for this particular use case as I could have used data frame-based tools and saved the result in a column.²⁹ However, there are limits to such workflows. For example, I recently trained a complex ML model for different variables and parameters, while also loading a different dataset for each combination. Loops are useful in such cases, and lists a natural vehicle to structure complex outputs. The main issue with nested lists is that they are complex to query. What if we wanted to know the R^2 of these 6 models? We would need to use, e.g., `resultsARGEMP$summary$r.squared` for each model.

This nested list-access problem was the main reason for creating `get_elem()`: an efficient recursive list-filtering function which, by default, simplifies the list tree as much as possible.

```
R> str(r_sq <- results |> get_elem("r.squared"))
```

```
List of 3
 $ ARG:List of 2
  ..$ EMP: num 0.907
  ..$ VA : num 1
 $ BRA:List of 2
  ..$ EMP: num 0.789
  ..$ VA : num 0.999
 $ CHL:List of 2
  ..$ EMP: num 0.106
  ..$ VA : num 0.999
```

²⁹E.g., `GGDC10S |> fgroup_by(Country, Variable) |> fsummarise(results = my_fun(lm(log(AGR+1) log(MIN+1) + log(MAN+1) + Year)))` with `my_fun <- function(m) list(list(m, BIC(m), summary(m)))`.


```
R> rowbind(r_sq, idcol = "Country", return = "data.frame")
```

	Country	EMP	VA
1	ARG	0.9068	0.9996
2	BRA	0.7888	0.9988
3	CHL	0.1058	0.9991

Note how the "summary" branch was eliminated since it is common to all final nodes; `results |> get_elem("r.squared", keep.tree = TRUE)` could have been used to keep it. `rowbind()` then efficiently combines lists of lists. We can also apply `t_list()` to turn the list inside-out:

```
R> r_sq |> t_list() |> rowbind(idcol = "Variable", return = "data.frame")
```

	Variable	ARG	BRA	CHL
1	EMP	0.9068	0.7888	0.1058
2	VA	0.9996	0.9988	0.9991

`rowbind()` is limited if `get_elem()` returns a more nested or asymmetric list, potentially with vectors/arrays in the final nodes. Suppose we wanted to extract the coefficient matrices:

```
R> results$ARG$EMP$summary$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	26.583617	1.2832583	20.7157	1.747e-28
log(MIN + 1)	0.083168	0.0352493	2.3594	2.169e-02
log(MAN + 1)	-0.064413	0.0767614	-0.8391	4.048e-01
Year	-0.009683	0.0005556	-17.4278	1.003e-24

For such cases, `unlist2d()` provides a complete recursive generalization of `unlist()`. It creates a 'data.frame' or 'data.table' representation of any nested list using recursive row-binding and coercion operations, while generating (optional) id variables representing the list tree and (optionally) saving row names of matrices or data frames. In the present example

```
R> results |> get_elem("coefficients") |> get_elem(is.matrix) |>
+   unlist2d(idcols = c("Country", "Variable"),
+   row.names = "Covariate") |> head(3)
```

	Country	Variable	Covariate	Estimate	Std. Error	t value	Pr(> t)
1	ARG	EMP	(Intercept)	26.58362	1.28326	20.7157	1.747e-28
2	ARG	EMP	log(MIN + 1)	0.08317	0.03525	2.3594	2.169e-02
3	ARG	EMP	log(MAN + 1)	-0.06441	0.07676	-0.8391	4.048e-01

where `get_elem(is.matrix)` is needed because the models also contain "coefficients".

This exemplifies the power of these tools to create, query, and combine nested data in very general ways. Further useful functions include `has_elem()` to check for the existence of

elements, `ldepth()` to return the maximum level of recursion, `is_unlistable()` to check whether a list has atomic elements in all final nodes, `[ir]reg_elem()` to recursively extract the (non-)atomic elements, and `rapply2d()` to apply functions to nested lists of data objects.

7. Summary statistics

collapse's **summary statistics functions** offer a parsimonious yet powerful toolkit to examine complex datasets. A particular focus has been on providing tools to understand multilevel (panel) data. Recall the World Development panel dataset (**wlddev**) from Section 3. The function `varying()` can be used to examine which of the variables are time-varying.

```
R> varying(wlddev, ~ iso3c)
```

country	date	year	decade	region
FALSE	TRUE	TRUE	TRUE	FALSE
income	OECD	PCGDP	LIFEEX	GINI
FALSE	FALSE	TRUE	TRUE	TRUE
ODA	POP	center_PCGDP	center_LIFEEX	
TRUE	TRUE	TRUE	TRUE	

A related exercise is to decompose the variance of a panel series into variation between countries and variation within countries over time. Using the (de-)meaning functions `fbetween()`/`fwithin()` supporting 'indexed_series' (Table 2), this is easily demonstrated:

```
R> LIFEEXi <- reindex(wlddev$LIFEEX, index = wlddev$iso3c)
R> all.equal(fvar(LIFEEXi), fvar(fbetween(LIFEEXi)) + fvar(fwithin(LIFEEXi)))

[1] TRUE
```

The `qsu()` (quick-summary) function provides an efficient method to compute this decomposition, considering the group-means instead of the between transformation³⁰ and adding the overall mean back to the within transformation to preserve the scale of the data.

```
R> qsu(LIFEEXi)
```

	N/T	Mean	SD	Min	Max
Overall	11670	64.2963	11.4764	18.907	85.4171
Between	207	64.9537	9.8936	40.9663	85.4171
Within	56.3768	64.2963	6.0842	32.9068	84.4198

The decomposition above implies more variation in life expectancy between countries than within countries over time. It can also be computed for different subgroups and with sampling weights. `qsu()` also has a data frame method, and by default computes simple statistics.³¹ Below, I take the latest post-2015 estimates and summarise LIFEEX by income groups with population weights. The **WeightSum** column thus records the total population in each group.

³⁰This is more efficient and equal to using the between transformation if the panel is balanced.

³¹The `pid` argument to `qsu()` can also be used to manually pass identifiers for panel-decomposition, e.g., `pid = iso3c`. With indexed data, it is automatically set to the first column in the index (`effect = 1`).

```
R> wlda15 <- wlddev |> fsubset(year >= 2015) |> fgroup_by(iso3c) |> flast()
R> qsu(wlda15, by = LIFEEX ~ income, w = ~ POP)
```

	N	WeightSum	Mean	SD	Min	Max
High income	68	1.19122607e+09	80.879	2.441	70.6224	85.078
Low income	29	694'893773	63.8061	3.9266	53.283	72.697
Lower middle income	47	3.06353648e+09	68.7599	4.7055	54.331	76.699
Upper middle income	55	2.67050662e+09	75.9476	2.3895	58.735	80.279

For greater detail, `descr()` provides a rich (grouped, weighted) statistical description. It does not support panel-variance decompositions like `qsu()`, but also computes detailed frequency tables for categorical data. Below, I summarize income and LIFEEX by OECD membership, scaling the weights to limit printout and replacing missing weights with 0 (the default).³²

```
R> descr(wlda15, by = income + LIFEEX ~ OECD, w = ~ replace_na(POP / 1e6))
```

Dataset: wlda15, 2 Variables, N = 216, WeightSum = 7620.902563

Grouped by: OECD [2]

	N	Perc	WeightSum	Perc
FALSE	180	83.33	6311.15	82.81
TRUE	36	16.67	1309.75	17.19

income (factor): Income Level

Statistics (WeightSum = 7621, 0% NAs)

	WeightSum	Perc	Ndist
FALSE	6311.15	82.81	4
TRUE	1309.75	17.19	2

Table (WeightSum Perc)

	FALSE		TRUE		Total	
Lower middle income	3064	48.5	0	0.0	3064	40.2
Upper middle income	2460	39.0	211	16.1	2671	35.0
High income	93	1.5	1099	83.9	1192	15.6
Low income	695	11.0	0	0.0	695	9.1

LIFEEX (numeric): Life expectancy at birth, total (years)

Statistics (N = 199, 7.87% NAs)

	N	Ndist	WeightSum	Perc	Mean	SD	Min	Max	Skew	Kurt
FALSE	163	164	6310.41	82.81	71.14	5.77	53.28	85.08	-0.99	3.76
TRUE	36	36	1309.75	17.19	80.32	2.77	75.05	84.36	-0.29	2.11

Quantiles

	1%	5%	10%	25%	50%	75%	90%	95%	99%
FALSE	54.69	59.38	63.67	69.64	71.71	76.89	76.91	76.91	77.94
TRUE	75.07	75.14	76.03	78.77	80.86	82.86	83.61	84.05	84.3

³²This is necessary in `descr()` because `fquantile()` does not support missing weights for non-missing x.

`descr()` also has a `stepwise` argument to describe one variable at a time, allowing users to naturally ‘step-through’ the variables in a large dataset while spreading the computational burden. The [documentation](#) provides more details and examples. Both `qsu()` and `descr()` come with an `as.data.frame()` method for efficient tidying and easy further analysis.

A final noteworthy function from **collapse**’s descriptive statistics toolkit is `qtab()`, an enhanced drop-in replacement for `table()`. It is enhanced both in a statistical and a computational sense, providing a remarkable performance boost, an option `sort = FALSE` to preserve the first-appearance-order of vectors being cross-tabulated, support for frequency weights (`w`), and the ability to compute different statistics representing table entries using these weights—vectorized when using *Fast Statistical Functions* as demonstrated below.

```
R> wlda15 /> with(qtab(OECD, income))
```

	income			
OECD	High income	Low income	Lower middle income	Upper middle income
FALSE	45	30	47	58
TRUE	34	0	0	2

This shows the total population (latest post-2015 estimates) in millions.

```
R> wlda15 /> with(qtab(OECD, income, w = POP / 1e6))
```

	income			
OECD	High income	Low income	Lower middle income	Upper middle income
FALSE	93.01	694.89	3063.54	2459.71
TRUE	1098.75	0.00	0.00	211.01

This shows the average life expectancy in years. The use of `fmean()` toggles an efficient vectorized computation of the table entries (i.e., `fmean()` is only called once).

```
R> wlda15 /> with(qtab(OECD, income, w = LIFEEX, wFUN = fmean))
```

	income			
OECD	High income	Low income	Lower middle income	Upper middle income
FALSE	78.75	62.81	68.30	73.81
TRUE	81.09			76.37

Finally, this calculates a population-weighted average of life expectancy in each group.

```
R> wlda15 /> with(qtab(OECD, income, w = LIFEEX, wFUN = fmean,
+                      wFUN.args = list(w = POP)))
```

	income			
OECD	High income	Low income	Lower middle income	Upper middle income
FALSE	77.91	63.81	68.76	75.93
TRUE	81.13			76.10

‘qtab’ objects inherit the ‘table’ class, thus all ‘table’ methods apply. Apart from the above summary functions, **collapse** also provides `pwcor()`, `pwcov()`, and `pwnobs()` for convenient (pairwise, weighted) correlations, covariances, and observations counts, respectively.

8. Global options

collapse is **globally configurable** to an extent few packages are: the default values of key function arguments governing the behavior of its algorithms, and the exported namespace, can be adjusted interactively through the `set_collapse()` function. These options are saved in an internal environment called `.op`. Its contents can be accessed using `get_collapse()`.

The current set of options comprises the default behavior for missing values (`na.rm` arguments in all statistical functions and algorithms), sorted grouping (`sort`), multithreading and algorithmic optimizations (`nthreads`, `stable.algo`), presentational settings (`stub`, `digits`, `verbose`), and, surpassing all else, the package namespace itself (`mask`, `remove`).

As evident from previous sections, **collapse** provides performance-improved or otherwise enhanced versions of functions already present in base R (like the *Fast Statistical Functions*, `funique()`, `fmatch()`, `fsubset()`, `ftransform()`, etc.) or other packages (especially **dplyr** (Wickham *et al.* 2023): `fselect()`, `fsummarise()`, `fmutate()`, `frename()`, etc.). The objective of being namespace compatible warrants such a naming convention, but this has a syntactical cost, particularly when **collapse** is used as the primary workhorse package.

To reduce this cost, **collapse**’s `mask` option allows masking existing R functions with the faster **collapse** versions by creating additional functions in the namespace and instantly exporting them. All **collapse** functions starting with an ‘f’ can be passed to the option (with or without the ‘f’), e.g., `set_collapse(mask = c("subset", "transform"))` creates `subset <- fsubset` and `transform <- ftransform` and exports them. Special functions are `"n"`, `"table"/"qtab"`, and `"%in%"`, which create `n <- GRPN` for use in `(f)summarise/(f)mutate`, `table <- qtab`, and replace `%in%` with a faster version based on `fmatch()`, respectively. There also exist several **convenience keywords to mask related groups of functions**, such as `"manip"` (only data manipulation functions), or `"all"` (everything), as demonstrated below.

```
set_collapse(mask = "all")
wlddev |> subset(year >= 1990 & is.finite(GINI)) |>
  group_by(year) |>
  summarise(n = n(), across(PCGDP:GINI, mean, w = POP))
with(mtcars, table(cyl, vs, am))
sum(mtcars)
diff(EuStockMarkets)
mean(num_vars(iris), g = iris$Species)
unique(wlddev, cols = c("iso3c", "year"))
range(wlddev$date)
wlddev |> index_by(iso3c, year) |>
  mutate(PCGDP_lag = lag(PCGDP),
         PCGDP_growth = growth(PCGDP)) |> unindex()
```

The above is now 100% **collapse** code. Similarly, using this option, all code in this article could have been written without `f`-prefixes. Thus, **collapse** is able to offer a fast and syntactically clean experience of R—without the need to even restart the session. Masking is completely and interactively reversible: calling `set_collapse(mask = NULL)` instantly removes the additional functions. Option `remove` can further be used to remove (un-export) any **collapse** function, allowing manual conflict management. Function `fastverse::fastverse_conflicts()` from the related **fastverse** project (Krantz 2024) can be used to display namespace conflicts with **collapse**. Invoking either `mask` or `remove` detaches **collapse** and re-attaches it at the top of the search path, letting its namespace to take precedence over other packages.

Such global powers confer responsibilities upon package developers, as further elucidated in the **vignette on developing with collapse**. As a general rule, options `mask` and `remove` should be off-limits inside packages, and other options need to be reset immediately using `on.exit()`.

9. Benchmark

This section provides several simple benchmarks to show that **collapse** provides best-in-R performance for statistics and data manipulation on moderately sized datasets. They are executed on a 2024 Apple MacBook Pro with 48GB M4 Pro chip. It also discusses results from **third-party benchmarks** involving **collapse**. The first set of benchmarks show that **collapse** provides faster computationally intensive operations like unique values and matching on large integer and character vectors. It creates integer/character vectors of 10 million obs, with 1000 unique integers and 5776 unique strings, respectively, which are deduplicated/matched in the benchmark. These fast basic operations impact many critical components of the package.

```
R> set.seed(101);
R> int <- 1:1000; g_int <- sample.int(1000, 1e7, replace = TRUE)
R> char <- c(letters, LETTERS, month.abb, month.name)
R> g_char <- sample(char <- outer(char, char, paste0), 1e7, TRUE)
R> bmark(base = unique(g_int), collapse = funique(g_int))
```

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	base	45.15ms	48.17ms	166.2MB	62	62	3.02s
2	collapse	5.52ms	7.74ms	38.2MB	371	75	3s

```
R> bmark(base = unique(g_char), collapse = funique(g_char))
```

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	base	71.9ms	73.4ms	166.2MB	41	41	3.02s
2	collapse	14.8ms	17.7ms	38.2MB	165	33	3.01s

```
R> bmark(base = match(g_int, int), collapse = fmatch(g_int, int))
```

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	base	17.1ms	17.58ms	76.3MB	79	78	1.4s
2	collapse	7.48ms	8.12ms	38.2MB	260	70	2.11s

```
R> bmark(base = match(g_char, char), data.table =
+       chmatch(g_char, char), collapse = fmatch(g_char, char))
```

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	base	58.6ms	60.6ms	114.5MB	49	49	3.02s
2	data.table	32.8ms	34.2ms	38.1MB	86	22	3.01s
3	collapse	19.4ms	20.3ms	38.1MB	140	29	3s

The second set below shows that **collapse**'s statistical functions are very efficient in aggregating a numeric matrix with 10,000 rows and 1000 columns. They are faster than base R even without multithreading, but using 4 threads in this case induces a sizeable difference.

```
R> set_collapse(na.rm = FALSE, sort = FALSE, nthreads = 4)
R> m <- matrix(rnorm(1e7), ncol = 1000)
R> bmark(base = colSums(m), collapse = fsum(m))
```

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	base	5.4ms	6.12ms	7.86KB	492	0	3.01s
2	collapse	292.5µs	387.94µs	7.86KB	7352	1	3s

```
R> bmark(base = colMeans(m), collapse = fmean(m))
```

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	base	5.69ms	6.02ms	32.09KB	499	0	3s
2	collapse	285.24µs	393.35µs	7.86KB	7200	0	3s

```
R> bmark(matrixStats = matrixStats::colMedians(m), collapse = fmedian(m))
```

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	matrixStats	77.1ms	78ms	89.9KB	38	1	2.98s
2	collapse	19.4ms	19.6ms	27.2KB	154	0	3.01s

Below, I also benchmark a grouped version summing the columns within 1000 random groups.

```
R> g <- sample.int(1e3, 1e4, TRUE)
R> bmark(base = rowsum(m, g), collapse = fsum(m, g))
```

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	base	5.38ms	5.63ms	7.87MB	466	35	2.63s
2	collapse	918.11µs	1.19ms	7.71MB	1714	127	2.2s

I now turn to basic operations on a medium-sized real-world database recording all flights from New York City (EWR, JFK, and LGA) in 2023—provided by the **nycflights23** package. The **flights** table has 435k flights, and grouping it by day and route yields 76k unique trips.

```
R> fastverse_extend(nycflights23, dplyr, data.table); setDTthreads(4)
R> list(flights, airports, airlines, planes, weather) |> supply(nrow)
```



```
[1] 435352    1255      14    4840   26207
```

```
R> flights |> fselect(month, day, origin, dest) |> fnunique()
```

```
[1] 75899
```

In the following, I select 6 numeric variables and sum them across the 76k trips using **dplyr**, **data.table**, and **collapse**. Ostensibly, despite `sum()` being 'primitive' (implemented in C), there is a factor 100 between **dplyr**'s split-apply-combine and **collapse**'s fully vectorized execution.

```
R> vars <- .c(dep_delay, arr_delay, air_time, distance, hour, minute)
R> bmark(dplyr = flights |> group_by(month, day, origin, dest) |>
+       summarise(across(all_of(vars), sum), .groups = "drop"),
+       data.table = qDT(flights)[, lapply(.SD, sum), .SDcols = vars,
+       by = .(month, day, origin, dest)],
+       collapse = flights |> fgroup_by(month, day, origin, dest) |>
+       get_vars(vars) |> fsum())
```

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	dplyr	356.33ms	490.66ms	50.5MB	7	69	3.24s
2	data.table	7.37ms	8.5ms	20.8MB	315	33	3s
3	collapse	3.21ms	3.78ms	9.2MB	696	27	3s

Below, I also benchmark the mean and median functions in the same way. It is evident that with non-primitive R functions the split-apply-combine logic becomes even more costly.

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	dplyr_mean	1.14s	1.25s	48.57MB	3	69	3.69s
2	data.table_mean	7.5ms	8.69ms	18.93MB	314	32	3.01s
3	collapse_mean	3.27ms	3.93ms	9.11MB	643	27	3s

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	dplyr_median	4.21s	4.21s	52.8MB	1	84	4.21s
2	data.table_median	21.92ms	23.71ms	18.9MB	122	13	3s
3	collapse_median	9.37ms	10.57ms	11.1MB	255	13	3.01s

So far, **data.table**, by virtue of its internal vectorizations (also via dedicated grouped C implementations of simple functions), is competitive.³³ Below, I compute the range of one column (`x`) using `max(x) - min(x)`. As elucidated in Section 3, this expression is also vectorized in **collapse**, where it amounts to `fmax(x, g) - fmin(x, g)`, but not in **data.table**.

```
R> bmark(dplyr = flights |> group_by(month, day, origin, dest) |>
+       summarise(rng = max(arr_delay) - min(arr_delay), .groups = "drop"),
+       data.table = qDT(flights)[, .(rng = max(arr_delay) - min(arr_delay)),
+       by = .(month, day, origin, dest)],
+       collapse = flights |> fgroup_by(month, day, origin, dest) |>
+       fsummarise(rng = fmax(arr_delay) - fmin(arr_delay)))
```

³³Much longer data will likely also favor **data.table** over **collapse** due to its sub-column-level parallel grouping and implementation of simple functions like `sum()` and `mean()`, see, e.g., the [DuckDB Benchmarks](#).

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	dplyr	91.78ms	107.7ms	20.18MB	27	55	3.14s
2	data.table	55.02ms	67.2ms	5.77MB	45	27	3.05s
3	collapse	5.17ms	5.7ms	6.8MB	491	14	3.01s

I also benchmark table joins and pivots. The following demonstrates how all tables can be joined together using **collapse** and its default first-match left-join, which preserves **flights**.

```
R> flights |> join(weather, on = c("origin", "time_hour")) |>
+   join(planes, on = "tailnum") |> join(airports, on = c(dest = "faa")) |>
+   join(airlines, on = "carrier") |> dim()
```

```
left join: flights[origin, time_hour] 434526/435352 (99.8%) <21.94:1st> weat
duplicate columns: year, month, day, hour => renamed using suffix '_weather'
left join: x[tailnum] 424068/435352 (97.4%) <87.62:1st> planes[tailnum] 4840
duplicate columns: year => renamed using suffix '_planes' for y
left join: x[dest] 435352/435352 (100%) <3689.42:1st> airports[faa] 118/1255
left join: x[carrier] 435352/435352 (100%) <31096.57:1st> airlines[carrier]
duplicate columns: name => renamed using suffix '_airlines' for y
[1] 435352      48
```

The verbosity of `join()` is essential to understanding what has happened here—how many records from each table were matched and which duplicate non-id columns were suffixed with the (default) y-table name. Usually, I would set `drop.dup.cols = "y"` as keeping them is not helpful in this case, but the other packages don't have this option. For the benchmark, I set `verbose = 0` in **collapse** and employ the fastest syntax for **dplyr** and **data.table**:³⁴

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	dplyr_joins	213.13ms	265ms	559.4MB	12	58	3.14s
2	data.table_joins	173.31ms	229ms	491MB	14	62	3.14s
3	collapse_joins	9.94ms	14ms	89.7MB	182	100	3.05s

Evidently, the vectorized hash join provided by **collapse** is 10x faster than **data.table** on this database, at a substantially lower memory footprint. It remains competitive on **big data**.³⁵

Last but not least, I benchmark pivots, starting with a long pivot that simply melts the 6 columns aggregated beforehand into one column, duplicating all other columns 6 times:

```
R> bmark(tidyr = tidyr::pivot_longer(flights, cols = vars),
+       data.table = qDT(flights) |> melt(measure = vars),
+       collapse = pivot(flights, values = vars))
```

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	tidyr	72.4ms	134ms	254MB	25	54	3.12s
2	data.table	43ms	59ms	209MB	45	34	3.11s
3	collapse	14.1ms	25ms	209MB	77	91	3.02s

³⁴`left_join(..., multiple = "first")` for **dplyr** and `y[x, on = ids, mult = "first"]` for **data.table**.

³⁵**data.table** joins utilize multithreaded radix-ordering—a very different logic more useful for big data.

Memory-wise, **collapse** and **data.table** are equally efficient, but **collapse** is faster, presumably due to more extensive use of `memset()` to copy values in C or smaller R-level overheads.

To complete the picture, I also perform a wide pivot where the 6 columns are summed (for efficiency) across the 3 origin airports and expanded to create 18 airport-value columns.

```
R> bmark(tidyr = tidyr::pivot_wider(flights, id_cols = .c(month, day, dest),
+   names_from = "origin", values_from = vars, values_fn = sum),
+   data.table = dcast(qDT(flights), month + day + dest ~ origin,
+   value.var = vars, fun = sum),
+   collapse_fsum = pivot(flights, .c(month, day, dest), vars,
+   "origin", how = "wider", FUN = fsum),
+   collapse_itnl = pivot(flights, .c(month, day, dest), vars,
+   "origin", how = "wider", FUN = "sum"))
```

	expression	min	median	mem_alloc	n_itr	n_gc	total_time
1	tidyr	365.57ms	404.26ms	143MB	8	62	3.25s
2	data.table	222.76ms	231.33ms	21.7MB	13	44	3.08s
3	collapse_fsum	6.31ms	8.45ms	39.1MB	269	68	3s
4	collapse_itnl	4.03ms	5.19ms	12.4MB	513	38	3s

Again, **collapse** is fastest, as it offers full vectorization, either via `fsum()`, which translates to `fsum(x, g, TRA = "fill")` before pivoting and thus entails a full deep copy of the `vars` columns, or via the optimized internal sum function which sums values 'on the fly' during the reshaping process. **data.table** is not vectorized here but at least memory efficient.

In summary, these benchmarks show that **collapse** provides outstanding performance and memory efficiency on a typical medium-sized real-world database popular in the R community.

9.1. Other benchmarks

The [DuckDB Benchmarks](#) compare many software packages for database-like operations using large datasets (big data) on a linux server. The January 2025 run distinguishes 6 packages that consistently achieve outstanding performances: **DuckDB**, **Polars**, **ClickHouse**, **Apache Datafusion**, **data.table**, and **collapse**. Of these, **DuckDB**, **ClickHouse**, and **Datafusion** are vectorized database (SQL) engines, and **Polars** is a Python/Rust DataFrame library and SQL engine. These four are supported by (semi-)commercial entities, leaving **data.table** as the only fully community-led project, and **collapse** as the only project that is single-authored and without financial support. The benchmarks show that **collapse** achieves the highest relative performance on 'smaller' (10-100 million row) datasets and performing advanced operations.

Since June 2024, there is also an independent [database-like operations benchmark](#) by [Adrian Antico](#) using a windows server and executing scripts inside IDEs (VScode, Rstudio), on which **collapse** achieved the overall fastest runtimes. I also very recently started a [user-contributed benchmark Wiki](#) as part of the [fastverse project](#) promoting high-performance software for R, where users can freely contribute benchmarks involving, but not limited to, **fastverse** packages. These benchmarks agree that **collapse** offers a computationally outstanding experience, particularly for medium-sized datasets, complex tasks, and on users PCs/Macs—which typically have smaller memory and parallel computing resources but faster chips than servers.

9.2. Limitations and outlook

collapse maximizes three principal objectives: being class-agnostic/fully compatible with the R ecosystem (supporting statistical operations on vector, matrix and data.frame-like objects), being statistically advanced, and being fast. This warranted some design choices away from maximum performance for large data manipulation.³⁶ Its limited use of multithreading and SIMD instructions, partly by design constraints and by R's C API, and the use of standard types for internal indexing, imposes hard-limits—the maximum integer in R is 2,147,483,647 → the maximum vector length **collapse** supports. It is and will remain an in-memory tool.

Despite these constraints, **collapse** provides very respectable performance even on very large datasets by virtue of its algorithmic and memory efficiency. It is, together with the popular **data.table** package offering more sub-column-level parallel architecture for basic operations, well-positioned to remain a premier tool for in-memory statistics and data manipulation.

10. Conclusion

collapse was first released to CRAN in March 2020, and has grown and matured considerably over the course of 5 years. It has become a new foundation package for statistical computing and data transformation in R—one that is statistically advanced, class-agnostic, flexible, fast, lightweight, stable, and able to manipulate complex scientific data with ease. As such, it opens up new possibilities for statistics, research, production, and package development in R.

This article provides a quick guide to the package, articulating its key ideas and design principles and demonstrating all core features. At this point the API is stable—it has changed very little over the 5 years and no further changes are planned. Compatibility with R version 3.5.0 will be maintained for as long as possible. Minor new features are currently planned.

For deeper engagement with **collapse**, visit its [website](#) or start with the vignette summarizing all available [documentation and resources](#). Users can also follow **collapse** on [Twitter/X](#) and [Bluesky](#) to be notified about major updates and participate in community discussions.

Finally, **collapse** users are also encouraged to familiarize themselves with the **fastverse**,³⁷ a suite of complementary high-performance packages for statistical computing and data manipulation in R that offer more advanced tools in several statistical computing domains. The **fastverse** metapackage additionally provides a lightweight framework to jointly load and manage these packages, as well as to build customized and fully separate package verses.

³⁶Which nowadays would demand creating a multithreaded, vectorized query engine with optimized memory buffers/vector types to take full advantage of SIMD processing as in **DuckDB** or **Polars**. Such an architecture is very difficult to square with R vectors and R's 30-year old C API.

³⁷Website: <https://fastverse.github.io/fastverse/>

Computational details

The results in this paper were obtained using R (R Core Team 2025) 4.4.3 with **collapse** 2.1.2, **data.table** 1.17.0, **dplyr** 1.1.4, **tidyr** 1.3.1, **matrixStats** 1.5.0, **fastverse** 0.3.4, **nycflights23** 0.2.0 (Ismay, Couch, and Wickham 2024), and **bench** 1.1.4 (Hester and Vaughan 2025). All packages used are available from the Comprehensive R Archive Network (CRAN) at <https://CRAN.R-project.org/>. The benchmark was run on an Apple MacBook Pro (2024) with a 48GB M4 Pro processor (single core speed ~ 4.4 GHz). Packages were compiled from source using Apple Clang version 17.0.0 with OpenMP enabled and the -O2 flag.

The `bmark()` function used for benchmarking is defined as follows:

```
bmark <- function(...) {  
  bench::mark(..., min_time = 3, check = FALSE) |>  
    janitor::clean_names() |>  
    fselect(expression, min, median, mem_alloc, n_itr, n_gc, total_time) |>  
    fmutate(expression = names(expression)) |>  
    dapply(as.character) |> qDF()  
}
```

Acknowledgments

The source code of **collapse** has been heavily inspired by (and partly copied from) **data.table** (Matt Dowle and Arun Srinivasan), R's source code (R Core Team and contributors worldwide), the **kit** package (Morgan Jacob), and **Rcpp** (Dirk Eddelbuettel). Packages **plm** (Yves Croissant, Giovanni Millo, and Kevin Tappe) and **fixest** (Laurent Berge) have also provided a lot of inspiration (and a port to its demeaning algorithm in the case of **fixest**). I also thank many people from diverse fields for helpful answers on Stackoverflow and many other people for encouragement, feature requests, and helpful issues and suggestions.

References

- Bengtsson H (2025). **matrixStats**: *Functions that Apply to Rows and Columns of Matrices (and to Vectors)*. R package version 1.5.0, URL <https://CRAN.R-project.org/package=matrixStats>.
- Bergé L (2018). “Efficient Estimation of Maximum Likelihood Models with Multiple Fixed-Effects: the R Package **FENmlm**.” *CREA Discussion Papers*, (13).
- Bouchet-Valat M, Kamiński B (2023). “**DataFrames.jl**: Flexible and Fast Tabular Data in Julia.” *Journal of Statistical Software*, **107**(4), 1–32. doi:10.18637/jss.v107.i04. URL <https://www.jstatsoft.org/index.php/jss/article/view/v107i04>.
- Chau J (2022). **rrapply**: *Revisiting Base Rapply*. R package version 1.2.6, URL <https://CRAN.R-project.org/package=rrapply>.
- Croissant Y, Millo G (2008). “Panel Data Econometrics in R: The **plm** Package.” *Journal of Statistical Software*, **27**(2), 1–43. doi:10.18637/jss.v027.i02.
- Dowle M, Srinivasan A (2025). **data.table**: *Extension of ‘data.frame’*. R package version 1.17.0, URL <https://CRAN.R-project.org/package=data.table>.
- Harris CR, *et al.* (2020). “Array programming with **NumPy**.” *Nature*, **585**, 357–362. doi:10.1038/s41586-020-2649-2.
- Hester J, Vaughan D (2025). **bench**: *High Precision Timing of R Expressions*. R package version 1.1.4, URL <https://CRAN.R-project.org/package=bench>.
- Hyndman RJ, Fan Y (1996). “Sample Quantiles in Statistical Packages.” *American Statistician*, pp. 361–365.
- Ismay C, Couch SP, Wickham H (2024). **nycflights23**: *Flights and Other Useful Metadata for NYC Outbound Flights in 2023*. R package version 0.2.0, URL <https://CRAN.R-project.org/package=nycflights23>.
- JuliaStats (2023). “**StatsBase.jl**: Julia Package for Basic Statistics.” URL <https://github.com/JuliaStats/StatsBase.jl>.
- Krantz S (2024). **fastverse**: *A Suite of High-Performance Packages for Statistics and Data Manipulation*. R package version 0.3.4, URL <https://fastverse.github.io/fastverse/>.
- Larmarange J (2023). **labelled**: *Manipulating Labelled Data*. R package version 2.12.0, URL <https://CRAN.R-project.org/package=labelled>.
- Lumley T (2004). “Analysis of Complex Survey Samples.” *Journal of Statistical Software*, **9**(1), 1–19. R package version 2.2.
- McKinney W (2010). “Data Structures for Statistical Computing in Python.” In Stéfan van der Walt, Jarrod Millman (eds.), *Proceedings of the 9th Python in Science Conference*, pp. 56 – 61. doi:10.25080/Majora-92bf1922-00a.

- Müller K, Wickham H (2023). **tibble**: Simple Data Frames. R package version 3.2.1, URL <https://CRAN.R-project.org/package=tibble>.
- pandas Development Team (2023). “pandas-dev/pandas: **pandas**.” doi:10.5281/zenodo.10426137. URL <https://doi.org/10.5281/zenodo.10426137>.
- Papadakis M, et al. (2023). **Rfast**: A Collection of Efficient and Extremely Fast R Functions. R package version 2.1.0, URL <https://CRAN.R-project.org/package=Rfast>.
- Pebesma E (2018). “Simple Features for R: Standardized Support for Spatial Vector Data.” *The R Journal*, **10**(1), 439–446. doi:10.32614/RJ-2018-009. URL <https://doi.org/10.32614/RJ-2018-009>.
- R Core Team (2025). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. URL <https://www.R-project.org/>.
- Ryan JA, Ulrich JM (2023). **xts**: eXtensible Time Series. R package version 0.13.1, URL <https://CRAN.R-project.org/package=xts>.
- Signorell A (2023). **DescTools**: Tools for Descriptive Statistics. R package version 0.99.52, URL <https://CRAN.R-project.org/package=DescTools>.
- StataCorp LLC (2023). *STATA Statistical Software: Release 18*. College Station, TX. URL <https://www.stata.com>.
- Vink R, et al. (2023). “pola-rs/polars: Python **polars** 0.20.2.” doi:10.5281/zenodo.10413093. URL <https://doi.org/10.5281/zenodo.10413093>.
- Wang E, et al. (2020). “A New Tidy Data Structure to Support Exploration and Modeling of Temporal Data.” *Journal of Computational and Graphical Statistics*, **29**(3), 466–478. doi:10.1080/10618600.2019.1695624. URL <https://doi.org/10.1080/10618600.2019.1695624>.
- Wickham H (2007). “Reshaping Data with the **reshape** Package.” *Journal of Statistical Software*, **21**(12), 1–20. URL <http://www.jstatsoft.org/v21/i12/>.
- Wickham H (2014). “Tidy Data.” *Journal of Statistical Software*, **59**(10), 1–23. doi:10.18637/jss.v059.i10. URL <https://www.jstatsoft.org/index.php/jss/article/view/v059i10>.
- Wickham H, Henry L (2023). **purrr**: Functional Programming Tools. R package version 1.0.1, URL <https://CRAN.R-project.org/package=purrr>.
- Wickham H, et al. (2019). “Welcome to the **tidyverse**.” *Journal of Open Source Software*, **4**(43), 1686. doi:10.21105/joss.01686.
- Wickham H, et al. (2023). **dplyr**: A Grammar of Data Manipulation. R package version 1.1.4, URL <https://CRAN.R-project.org/package=dplyr>.
- Wickham H, et al. (2024). **tidyr**: Tidy Messy Data. R package version 1.3.1, URL <https://CRAN.R-project.org/package=tidyr>.

Affiliation:

Sebastian Krantz
Kiel Institute for the World Economy
Haus Welt-Club
Düsternbrooker Weg 148
24105 Kiel, Germany
E-mail: sebastian.krantz@ifw-kiel.de